

Diffusion Curvature for Fast, Point-wise, Noise-Resistant Geometric Featurization of Graphs and Pointclouds

Kincaid MacDonald¹, Dhananjay Bhaskar², Kaly Zhang², Ian Adelstein³,
Smita Krishnaswamy^{4,5}

¹Yale,

²MILA,

³Yale Department of Math,

⁴Yale Department of Applied Math,

⁵Yale School of Medicine,

Corresponding author: Kincaid MacDonald, kincaid@aya.yale.edu

Abstract

For a number of years now work has been proceeding in order to bring to perfection the crudely conceived idea of a machine that would not only supply inverse reactive current for use in unilateral phase detractors, but would also be capable of automatically synchronizing cardinal grammeters. Such a machine is the “Turbo-Encabulator.”

Plain Language Summary

We introduce Diffusion Curvature, a fast, differentiable, noise-robust pointwise curvature for graphs and point clouds.

1 Introduction

One of the most ubiquitous subjects of analysis in data science is the humble point cloud. The points, by themselves, are high dimensional and noisy; it is up to the data scientist to wring sense out of them. Per the Manifold Hypothesis, we assume the points were sampled on or near the surface of a low-dimensional manifold embedded in high-dimensional Euclidean space. Manifold learning methods, like t-SNE, PHATE, and Diffusion Maps [(maaten2008VisualizingDataUsing?)](moon2019VisualizingStructureTrans endeavor to recover salient features of the underlying manifold, like geodesic distances, population clusterings, and dimension, from its high-dimensional noisy sampling.

Curvature is a particularly troublesome geometric property to translate into the discrete, sampled realm. In smooth Riemannian manifolds, curvature is a *local* phenomenon. It can be obtained by fitting osculating circles of radius limiting to zero, or computed from the manifold’s Hessian, using the Second Fundamental Form. None of these translate into the discrete realm. In a sampled manifold, taking a local limit is impossible – one can’t “zoom in” past the sampling of points – and we don’t have access to the parametrization of the manifold or its tangent bundle, without computationally costly and potentially error-prone estimation. Moreover, as our sampling is likely noisy, the curvature can only be recovered over a sufficiently large neighborhood of points to counter the spurious geometric artifacts created noisy sampling. Thus, in the discrete realm, curvature becomes a “semi-local” phenomenon, in which neither the smallest nor larger scales can be trusted.

There are elegant generalizations of classical curvature to discrete spaces that overcome many of these roadblocks. Ollivier’s *Coarse Ricci Curvature* (CRC) employs optimal transport theory to relate the behavior of a discrete neighborhood to its smooth counterpart (ollivier2009RicciCurvatureMarkov?). Sturm’s *displacement convexity of entropy* (DCE) measures the proliferation of midpoints in positive curvature (sturm2006GeometryMetricMeasure?). Both methods use optimal transport as the basis of their “semi-local” measurement. Rather than trying to zoom in on a point, they define curvature between pairs of points, approximating, at a coarse scale, a Ricci tangent vector.

Although these techniques are theoretically elegant, general, and applicable to any metric measure space, the setting of noisily sampled point clouds is practically challenging for CRC and DCE. Both methods rely on the graph’s shortest-path lengths as an approximation of the manifold’s ground distance - a perilous assumption when dealing with noisy data. And for large datasets, optimal transport calculations can be computationally prohibitive.

In this paper, we develop Diffusion Curvature, a fast curvature estimate derived solely from the graph diffusion matrix. We first introduced the ideas behind Diffusion Curvature in (bhaskarDiffusionCurvatureEstimating2022?), in which we demonstrated its ability to produce an unsigned magnitude of curvature estimation

for toy datasets and single-cell data, and proved a correspondence between the ratios of scalar curvature and diffusion curvature. We now present a refined definition which produces *signed* curvature values and prove bounds relating Diffusion Curvature to coarse Ricci Curvature. We demonstrate Diffusion Curvature’s robustness to noise and sampling artifacts, and position our technique as an adaptation of coarse Ricci curvature particularly suitable for point cloud data. ## Background

1.1 Curvature in the Continuous Setting

1.2 The Discrete Setting

Within the ambient setting of points $x_i \in \mathbb{R}^D$, the Euclidean distances between the points in our point cloud are not very useful. To perform geometric analysis, we want the manifold’s *geodesic* distances between $x_i, x_j \in \mathcal{M}$. However, manifolds are locally euclidean, so within a sufficiently small neighborhood of $x_i \in \mathcal{M}$, the euclidean distances are accurate. This is the basis of graph construction: retain only the trustworthy local distances, discard the rest, and then “integrate” over the local neighborhoods to recover features of the global geometry.

A graph $G = (V, E)$ is a collection of n vertices $v_i \in V$ connected by (possibly weighted) edges $e_{ij} \in E$. It is efficiently represented by a single *adjacency* (or *affinity*) matrix $A \in \mathbb{R}^{n \times n}$, where A_{ij} expresses the degree of connection between the vertices v_i and v_j . In a binary adjacency matrix, $A_{ij} = 1$ iff there is an edge between v_i and v_j . In a weighted affinity matrix, $0 < A_{ij} < 1$ with a higher affinity indicating a closer connection between the nodes.

One can construct an affinity matrix from a point cloud with the following algorithm: 1. Compute the matrix D of pairwise euclidean distances between points, so that $D_{ij} = \|x_i - x_j\|_2$. 2. Apply a kernel κ to the distances to construct the affinity matrix, where $A_{ij} = \kappa(D_{ij})$. This is typically the gaussian kernel:

$$k(y) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{y}{\sigma^2}\right)$$

There are a variety of heuristics for selecting an appropriate kernel bandwidth σ . In this paper, we use an adaptive kernel bandwidth, in which, when computing $k(D_{ij})$, σ is set to the mean distance from the points x_i and x_j to their k -th nearest neighbor.

After building our graph affinity matrix A , we created a new representation of the point cloud X – turning it from an $n \times D$ matrix of unwieldy ambient coordinates into an $n \times n$ matrix of pairwise connections between points. The challenge is now to reassemble this information of local connectivity to recover the features of \mathcal{M} . Graph diffusion does precisely this.

1.3 Graph Diffusion

The graph diffusion matrix P is a commonly-used method of “integrating” the local connectivity of the graph A into global geometric descriptors of \mathcal{M} . Coifman and Lafon (WAWA YEAR) proved a correspondence between iterated graph diffusion P^t and the Neumann heat kernel on \mathcal{M} . Their technique, *Diffusion Maps*, uses the euclidean distances between eigencoordinates of P to approximate the geodesic distances on \mathcal{M} . The visualization technique *PHATE* (moon2019VisualizingStructureTransitions?) constructs a low-dimensional embedding of a point cloud X such that a distance between the transition probabilities P of X is preserved in the embedding. (More on properties of phate, trajectory preservation.) *Diffusion Earth Mover’s Distance* (tongDiffusionEarthMover2021?) efficiently approximates the transportation distance between distributions on a graph using multi-scale wavelet transform obtained by applying different scales of diffusion. *LEGSNet*’s “learnable geometric scattering” computes tunable scales of diffusion with a graph neural network and achieves state of the art performance on biochemistry graph classification

(**tong2020DataDrivenLearningGeometric?**). These are but a few of the many manifold learning techniques based in diffusion.

Constructing the diffusion matrix from the affinity matrix A is straightforward: you simply row-normalize A , with an optional step to normalizing by density.

Here is the algorithm presented in Coifman and Lafon (**coifman2006DiffusionMaps?**):

1. (Optional) Compute an *anisotropic density normalization* on A , obtaining the anisotropic adjacency matrix A_* . 3. Construct the degree matrix D , whose diagonal entries are the rowsums of A , i.e. $D_{ii} = \sum_j A_{ij}$. The other entries are zeros. 4. Define $P = D^{-1}A$, the graph diffusion matrix.

□ Clean this up: get anisotropic equation, and clarify the role of the self affinity. When is it removed? When is laziness added?

P has several nice properties. The rows $P[i]$ give the transition probabilities of a single step random walk starting at point x_i ; each row $P[i]$ can be viewed as a probability distribution centered at x_i . This is preserved under powers of the matrix. The rows of P^t still sum to 1, and $P^t[i]$ now gives the probability distribution of a t -step random walk starting at x_i .

Although P is not symmetric, it is conjugate to a symmetric matrix, via $D^{0.5}PD^{-0.5} = D^{-0.5}AD^{-0.5}$, granting it a full basis of real-valued eigenvectors and eigenvalues. These eigenvectors are shared with the normalized graph Laplacian $L = I - D^{-0.5}AD^{-0.5}$. The eigenvalues of P lie between 0 and 1. Powering the matrix P^t thus corresponds to powering the eigenvalues λ_i^t of P , via diagonalization

$$P^t = \Psi \Lambda^t \Psi^T$$

This is similar to applying a low-pass filter to the graph. As t increases, the smallest eigenvalues decay fastest under repeated powering, and their corresponding eigenvector vanishes from the eigenbasis – leaving only the largest λ_i , whose eigenvectors trace global geometric features.

This is a remarkable feature of the diffusion matrix: the ability to “denoise” itself by iterating the random walk over larger time scales. Intuitively, the paths through the data most robustly trafficked by random walkers are those supported by multiple high-probability connections from independent starting points. ## Related Work

1.4 Foreman Ricci Curvature

1.5 Hickock’s Curvature

1.6 Ollivier-Ricci Curvature

Developed by Yann Ollivier in 2007, *Coarse Ricci Curvature* (or sometimes, “Ollivier Ricci Curvature”) is a direct translation of Ricci curvature to discrete metric spaces like graphs (**ollivier2009RicciCurvatureMarkov?**). Several classical properties of Ricci curvature can be extended to the graph setting using Coarse Ricci Curvature. Ollivier has, for instance, proven versions of concentration inequalities, Bonnet Myers (more). Coarse Ricci Curvature has, in this way, become something of a bridge between continuous and coarse geometry. The basis of this bridge is optimal transport, and specifically, the 1-Wasserstein distance.

In the Riemannian setting, Ricci curvature captures the phenomenon that, in positive curvature, “small spheres are closer (in transportation distance) than their centers are” (CITE 43 in ORC Paper). On the sphere, for instance, imagine two circles placed at the north and south poles: every point is closer to the corresponding point on the opposite pole than the centers. In negatively curved spaces, the discrepancy reverses, while in a flat space, the average distance between the points of the circles is the distance between the centers.

Coarse Ricci Curvature captures a similar phenomenon on graphs. Instead of spheres, it uses locally-centered probability distributions defined by random walks.

And to measure the distance between these walks, it uses the 1-Wasserstein (or Earth Mover’s) distance. We’ll briefly define each.

The 1-Wasserstein distance is a measure of the distance between probability distributions. Given distributions μ_x and μ_y over some shared space X , the Wasserstein distance quantifies the smallest amount of “work” needed to transform one distribution into another, by transporting probability “mass” between pairs of points over the ground metric $d(x, y)$:

[!Info] Definition The 1-Wasserstein distance between distributions μ_x and μ_y is

$$W_1(\mu_x, \mu_y) := \inf_{\xi \in \Pi(\mu_x, \mu_y)} \int \int d(x, y) d\xi(x, y)$$

where the “transportation plan” ξ is drawn from the space $\Pi(\mu_x, \mu_y)$ of joint probability distributions over $X \times X$ which project onto μ_x and μ_y .

In the discrete setting, this translates naturally into an infimum over a summation.

$$W_1(\mu_x, \mu_y) := \inf_{\xi \in \Pi(\mu_x, \mu_y)} \sum_{x \in X} \sum_{y \in X} d(x, y) \xi(x, y)$$

What is the analog on a graph of a “small sphere” around a point? Ollivier replaces spheres with a family of measures $\mu_x(\cdot)$ defined for each point x , where 1. Each $\mu_x(\cdot)$ depends measurably on x , i.e. the map $x \rightarrow \mu_x$ is measurable. 2. Each $\mu_x(\cdot)$ has finite first moment, or *Jump*, i.e. for some $o \in X$ $\int d(o, y) \mu_x(y) dx < \infty$.

The *Jump* $J(\mu_x)$ of a measure, a measure of its concentration around a central point, is a concept to which we’ll return.

$$J(\mu_x) = \int_{y \in X} d(x, y) \mu_x(y) dx$$

In graphs, Ollivier defines these μ_x as the probability distributions created by a single-step random walk from the point x . With a transition probability α , and equal probability of moving to each of x ’s neighbors on the graph, $\mu_x(x) = (1 - \alpha)$ and $\mu_x(y) = \alpha$ if $y \in N(x)$ or 0 otherwise. This is analogous to defining $m_x = Pe_x$, if P is the diffusion matrix created from a binary adjacency matrix. Note, however, that there is nothing limiting us to binary adjacency matrices, or even single steps of diffusion; the two conditions above are equally satisfied by weighted adjacency matrices and t -step diffusions, and in sparse or noisy graphs, this may be desirable.

Definition

The *Coarse Ricci Curvature* between x and y is

$$\kappa(x, y) := 1 - \frac{W_1(m_x, m_y)}{d(x, y)}$$

There are a number of provisos attached to this definition, which tries to approximate a continuous phenomenon within discrete constraints. These constraints, and the relationship between Ricci and Ollivier’s coarse Ricci curvature are illustrated Ollivier’s Example 2.6 (**ollivier2009RicciCurvatureMarkov?**):

Let (X, d) be a smooth Riemannian manifold of dimension d and let vol be the Riemannian volume measure. Let $\epsilon > 0$ small enough and consider the ball of radius ϵ around each point x . Let $x, y \in X$ be two sufficiently close points. Let v be the unit tangent vector at x directed towards y . The coarse Ricci curvature along v is then

$$\kappa(x, y) = \frac{\epsilon^2 \text{Ric}(v, v)}{2(d + 2)} + o(\epsilon^3 + \epsilon^2 d(x, y))$$

Hence the coarse Ricci curvature applied to a manifold recovers the Ricci curvature, up to a scaling factor contingent on dimension, and plus an error term that grows with the radius of ball and distance between points.

Ollivier’s choice not to scale $\kappa(x, y)$ by dimension is interesting, and likely motivated by his application of coarse Ricci curvature to graph-like spaces for which dimension isn’t clearly defined, like social networks. Within our domain of point-cloud data, incorporating dimension may be desirable; without it, spaces of high dimension can be conflated with spaces of low negative curvature but high dimension.

A result on coarse Ricci curvature which will prove useful concerns the *contraction (or expansion) of measure* that occurs under diffusion in spaces of positive (or negative) curvature.

Proposition 20 ([ollivier2009RicciCurvatureMarkov?](#))

Let (X, d, m) be a metric space with a random walk. Let $\kappa \in \mathbb{R}$. Then we have $\kappa(x, y) \geq \kappa$ for all $x, y \in X$ iff for any two probability distributions $\mu, \mu' \in \mathcal{P}(X)$ one has

$$W_1(\mu \star m, \mu' \star m) \leq (1 - k)W_1(\mu, \mu')$$

Where

$$\mu \star m := \int_{x \in X} d\mu(x) m_x dx$$

2 Theory & Methods

2.1 Definition

Given samples $X \subseteq M$ and a flattening map $\Phi : X \rightarrow \mathbb{R}^d$,

The t -step *Diffusion Curvature* of x is

$$k_t(x) = 1 - \frac{W_1(\delta_x, p_X^t(x))}{W_1(\delta_x, p_{\Phi(X)}^t(x))}$$

{eq-definition}

Where p_X^t is the t -step random walk over X , and $p_{\Phi(X)}^t$ is the same over the flattened points $\Phi(X)$. In both cases, the W_1 distance is taken with respect to manifold distances.